

Support QA Process — Overview

Why we're building this

Support ticket quality is currently reviewed inconsistently: Team Leads (TLs) vary in how strictly they score the same behaviors, so there's no reliable, standardized signal for coaching, consistency, or CSAT improvement. In 2026 the Support Team Lead group aligned on a shared QA rubric and is now building an internal, AI-assisted review process rather than buying a commercial QA platform outright.

Goal: The operational goal of the QA process is to use QA data to drive coaching, consistency across the team, and CSAT improvement, with SLA/first-response time tracked alongside quality so process breakdowns surface early.

We're using this QA process as the test case for Camunda's agentic orchestration capabilities: a closed Zendesk ticket is a bounded, structured input; the rubric provides clear evaluation criteria; and the split between mechanical checks, AI-assisted analysis, and human judgment provides a useful workflow for testing how deterministic steps, AI agents, enterprise data access, and human review can be orchestrated together.

The QA use case will instead be used to evaluate Camunda as infrastructure for agentic orchestration, governance and auditability, controlled enterprise data access, cost control, security and privacy, human oversight, and operational resilience. AI-generated QA assessments will be treated as provisional during the trial.

The rubric: 3 dimensions, 9 categories

Communication, Technical Advice, and Process Adherence, each broken into 3 categories. Every category is scored on a 4-point scale (4-Excellent / 3-Meets Standard / 2-Developing / 1-Does Not Meet). Explicitly not scored: raw CSAT, response length in isolation, or ticket status preference where both choices are defensible.

Dimension	Category	What “Meets Standard” looks like
Communication	Tone, Professionalism & Acknowledgment	Warm, professional, uses the customer's name; acknowledges their specific situation before presenting a solution; appropriate greeting and closing.
Communication	Next Steps & Ownership	Clearly defines the exact action, who's responsible, and when the customer can expect an update or resolution.
Communication	Call / Channel Judgment	Right channel chosen for the situation — a call was offered when the issue was complex, and call notes are captured.
Technical Advice	Root Cause Identified	Correctly identifies and explains the best-supported root cause based on the available evidence, not just the symptom
Technical Advice	Technical Accuracy	Accurate, complete, up-to-date guidance, tailored to the customer's level of understanding.
Technical Advice	Completeness (incl. Workaround vs. Fix)	Every part of the question answered; any workaround is clearly framed as temporary with resolution status communicated.
Process Adherence	Jira Bug Ticket	Bug documented promptly with a thorough Jira ticket, linked to the Zendesk case immediately.
Process Adherence	Internal Notes	Notes capture troubleshooting, decisions, next steps, and ownership throughout the ticket, with a closing note on resolution.
Process Adherence	Follow-Up Commitments	Commitments to the customer are honoured within the promised timeframe without them needing to chase.

Of the nine categories, Process Adherence lends itself most directly to automated or AI-assisted checks today: Jira linking is mechanically verifiable, and Internal Notes / Follow-Up Commitments can be checked for the presence of supporting evidence but still need human review for quality and judgment. Communication and Technical Advice remain human-owned decisions. During the trial, the AI may produce provisional assessments and supporting evidence for these categories to exercise Camunda's orchestration, data-access, governance, and human-review capabilities, but the trial is not intended to prove that these judgments can be reliably automated. Categories that do not apply to a ticket may be marked Not Applicable rather than being forced into the four-point scale.

Selecting tickets for manual review

The AI-assisted process runs against 100% of closed tickets, performing available mechanical checks and producing provisional, evidence-backed assessments against the applicable rubric categories. From that full population, a target 5–8% is routed to a human TL for manual review, drawn from four channels:

Channel	What it captures	Why it's in the mix
Low provisional AI scores	Every ticket scoring below “Meets Standard” on any category — reviewed in full, not sampled.	Highest coaching value; these are the tickets most likely to need intervention.
High-score spot-check	A random sample of top-scoring tickets.	Catches AI false positives and scoring drift before it goes unnoticed.
Random baseline	A random draw across all remaining scored tickets.	Keeps the sample representative, not just outlier-chasing.
Anomaly-flagged	Tickets AI flags as atypical: more back-and-forth than usual, signs of customer frustration/escalation, slower-than-usual response cadence.	Surfaces tickets that “feel off” even when the score looks fine.

TLs configure the mix: the sample percentage per channel, and the specific thresholds/signals that count as an anomaly (reply count, sentiment, response-time deviation) — so the selection logic can be tuned as priorities change, not fixed once and left alone.

Proposed logic: how it would work

1. Ticket closes in Zendesk → the (AI agents in) Camunda process (e.g. with process controller and controlled data access) retrieves a controlled ticket snapshot containing the thread, approved metadata, tags, and relevant identifiers.
2. Deterministic checks run for mechanically verifiable information, such as SLA timing, Jira linking, reply count, and response cadence.
3. An AI agent produces provisional assessments against the applicable rubric categories, attaching supporting evidence and a confidence level to judgment-based findings. (Where required, the agent may use a limited set of approved tools, such as retrieving the current rubric or looking up the Jira issue linked to the current Zendesk ticket - to test the controlled tool access)
4. A selection engine builds the human-review queue, normally targeting 5–8%, from four channels: low provisional scores, a spot-check sample of high scorers, a random baseline, and anomaly-flagged tickets.
5. TLs configure the selection engine — sample size per channel and what counts as an anomaly — so the mix can be tuned over time.
6. The TL reviews queued tickets, assigns the final Meets / Needs Improvement / Exceeds rating, and adds coaching notes. The TL remains the owner of the final QA decision.
7. The process records the provisional AI assessment, supporting evidence, tool activity, selection reason, final human rating, and any override for audit and trial evaluation.
8. Calibration loop: TLs periodically re-score the same tickets to resolve disagreements and refine the rubric, operating process, AI confidence thresholds, and anomaly criteria. During the trial, this loop will also be used to evaluate whether Camunda provides sufficient traceability and version control for process, model, prompt, rubric, and configuration changes.
9. Results roll up into a weekly Support Brief (avg. QA score, trends, a highlighted example) and a monthly Product/Dev Brief (bug trends, documentation gaps). During the trial, reporting should clearly distinguish provisional AI-generated results from human-confirmed QA results.

Trial evaluation focus

The primary purpose of the trial is to evaluate Camunda as infrastructure for enterprise agentic AI orchestration. The trial will focus on whether Camunda can:

- orchestrate deterministic workflow steps, AI agent activity, enterprise integrations, business rules, and human review in one traceable process;
- restrict agents to approved tools, data, records, and actions;
- provide sufficient governance, auditability, versioning, and human oversight;
- measure and control model calls, token use, tool calls, retries, processing time, and infrastructure costs;
- meet security, privacy, credential-management, data-retention, and access-control requirements;
- handle provider failures, connector failures, malformed responses, timeouts, budget limits, and manual fallbacks predictably;
- provide enough operational visibility for support and troubleshooting.

Agent-assessment quality is a secondary measure. The trial only requires the agent to produce structurally usable output often enough to exercise these infrastructure capabilities. It is not intended to demonstrate production-ready QA scoring.

Status: rubric and mechanical/judgment split are aligned across TLs; the operating SOP (ticket sampling method, review cadence, feedback capture) is still being finalized. That gap — turning an agreed rubric into a repeatable, partly-automated workflow — is exactly what we want to explore with AI orchestration.